



UNIVERSITÀ  
DEGLI STUDI  
DI MILANO



**MEC ITALY**

# **MONGODB PLAYGROUND**

**ALFIO FERRARA, STEFANO MONTANELLI**

Department of Computer Science

Via Comelico 39, 20135 Milano

{alfio.ferrara,stefano.montanelli}@unimi.it  
<http://islab.di.unimi.it/>

# SOFTWARE TOOLS

---

- **MongoDB server and client (v. 3.2.\*)**  
<https://www.mongodb.com/>
- **Python 2.7**  
<https://www.python.org/download/releases/2.7/>
- **Java SE 1.7 (or higher)**  
<https://www.oracle.com/java/index.html>
- **Python JSON libs**  
<https://docs.python.org/2/library/json.html>
- **Java JSON libs**  
<https://github.com/stleary/JSON-java>  
<https://github.com/google/gson>
- **Python MongoDB API (pymongo)**  
<https://api.mongodb.com/python/current/>
- **Java MongoDB API (java driver)**  
<http://mongodb.github.io/mongo-java-driver/3.2/>

# CHECK THE ENVIRONMENT

---

## MONGODB SERVER AND CLIENT

- **Laptop:** start server on localhost and then type **mongo** on your favourite shell console

## PYTHON

```
import pymongo
```

## JAVA

```
import org.json.*
```

# JSON IN A NUTSHELL

---

JSON means JavaScript Object Notation

Basic notions about syntax

- Data is organized in **key/value** pairs
- Data is separated by commas, objects are denoted by curly brackets, arrays are denoted by square brackets

```
{"game": "GTAV", "platform": "PS4", "genre": ["open-world", "shooting"], "price": 69.99}
```

# JSON VALUES (I)

---

A **value** corresponding to a key can be:

A number (integer or floating point)

```
{"price": 69.99, "version": 5}
```

A string (in double quotes)

```
{"game": "GTAV"}
```

A Boolean (true or false)

```
{"game": "GTAV", "in_collection": true}
```

## JSON VALUES (II)

---

A **value** corresponding to a key can be:

An array (in square brackets)

```
{"genre": ["open-world", "shooting"]}
```

An object (in curly braces)

```
{"game": "GTAV",  
  "owner": {  
    "first_name": "gino",  
    "last_name": "rossi"  
  }}
```

The null value

```
{"game": "GTAV", "year": null}
```

# NESTED OBJECTS

---

JSON objects can be nested at any level of depth

```
{ "game_shop": {  
  "name": "Gino's shop",  
  "address": {  
    "street": "via dei giochi",  
    "number": 16,  
    "city": "Milan",  
    "country": "Italy"  
  },  
  "games": [  
    { "game": "GTAV", "platform": "PS4",  
      "genre": [ "open-world", "shooting" ], "price": 69.99,  
      "reviews": [  
        { "evaluation": 5, "date": ISODate( "2015-06-29T00:00:00Z" ),  
          "reviewer": {  
            "first_name": "Mario", "last_name": "Bianchi"  
          }  
        },  
        { "evaluation": 4, "date": ISODate( "2016-07-16T00:00:00Z" ),  
          "reviewer": {  
            "first_name": "Maria", "last_name": "Verdi"  
          }  
        }  
      ]  
    }  
  ]  
}
```

# MISSING INFORMATION

Game

ID	name	platform	price
1	GTAV	PS4	69.99
2	Fallout4	PS4	NULL
3	Dying Light	PS4	NULL

```
{ "_id": 1, "name": "GTAV", "platform": "PS4", "price": 69.99 }
{ "_id": 2, "name": "Fallout4", "platform": "PS4", "price": null }
{ "_id": 3, "name": "Dying Light", "platform": "PS4" }

> db.games_example.find({"price": {"$gte": 40}})
{ "_id" : 1, "name" : "GTAV", "platform" : "PS4", "price" : 69.99 }

> db.games_example.find({"price": null})
{ "_id" : 2, "name" : "Fallout4", "platform" : "PS4", "price" : null }
{ "_id" : 3, "name" : "Dying Light", "platform" : "PS4" }

> db.games_example.find({"price": {"$exists": true}})
{ "_id" : 1, "name" : "GTAV", "platform" : "PS4", "price" : 69.99 }
{ "_id" : 2, "name" : "Fallout4", "platform" : "PS4", "price" : null }
```



## USE LIBS!

### PYTHON

```
import json
in_file = 'data/games.json'
with open(in_file, 'rU') as i_file:
    json.load(i_file, encoding='utf-8')
```

---

### JAVA

```
import org.json.*;
String inFile = "data/games.json";
try {
    String contents = new String(readAllBytes(get(inFile)));
    JSONObject data = new JSONObject(contents);
} catch (IOException e) {
    e.printStackTrace();
}
```

## USE LIBS AGAIN!

### PYTHON

```
import json
data = {'game': 'GTAV'}
out_file = 'data/games.json'
with open(out_file, 'wb') as out:
    json.dump(data, out, encoding='utf-8')
```

---

### JAVA

```
import org.json.*;
String outFile = "data/games.json";
JSONObject obj = new JSONObject();
obj.put("game", "GTAV");
try {
    FileWriter file = new FileWriter(outFile);
    file.write(obj.toJSONString());
} catch (IOException e) {
    e.printStackTrace();
}
```

## EXERCISE 1

---

- Download <http://islab.di.unimi.it/~alfio/mdbexe/games.json.gz>
- Read it in memory using your favourite programming language
- Understand data structure
- What kind of games? How many per kind?

## JSON VS BSON

---

MongoDB represents JSON documents in binary-encoded format called BSON behind the scenes. BSON extends the JSON model to provide additional data types and to be efficient for encoding and decoding within different languages.

<https://www.mongodb.com/json-and-bson>

# GO TO MONGODB

---

MongoDB documentation is awesome. So **use the docs!**

<https://docs.mongodb.com/manual/>

MongoDB is made of **databases** and **collections**. Forget CREATE DATABASE and CREATE COLLECTION stuff. If you need a DB and a collection **just start using it.**

## CONSOLE COMMANDS TO SURVIVE

```
> show dbs
> use [db_name]
> db.getCollectionNames()
> db.[collection_name].[function]
    (e.g., insert(), update(), remove(), find(), count(), aggregate())
> db.[collection_name].drop()
> db.dropDatabase()
```

# MONGODB BACKUP AND RESTORE

---

```
> mongodump --host [host] --db [db_name] --collection [collection_name] --out [path]
```

```
> mongorestore --collection [collection_name] --db [db_name] [path/file.bson]
```

```
> mongorestore [path/db_dir]
```

Use the mongo console to save your data and transfer to a different machine

# PUT DATA IN MONGODB PROGRAMMATICALLY

---

## PYTHON

```
import pymongo
db = pymongo.MongoClient(host_name)[db_name] #create client and select db
db[collection_name].insert_one({'key': 'value', ...}) #insert single record
db[collection_name].insert_many([{'key': 'value', ...}, {'key': 'value', ...}, ...])
                                #insert batch
```

## JAVA

```
import com.mongodb.*;
Mongo mongo = new Mongo(host_name, 27017);
DB db = mongo.getDB(db_name);
DBCollection a_collection = db.getCollection(collection_name);
JSONObject j = ...;
DBObject o = (DBObject) JSON.parse(j.toString()); //take care of mapping from JSON to BSON
a_collection.insert(o); //or even a List of DBObject for batch insert
```

## EXERCISE 2

---

Get data from games.json and put it in two collections, one for videogames and the other for boardgames

**Note** that the field 'date' in 'ratings' is encoded as a String in JSON but should be loaded as a Date in the DB

Check your results with respect to this dump: <http://islab.di.unimi.it/~alfio/mdbexe/ginfo.tgz>



# CRUD OPERATIONS AND QUERIES!

---

<https://docs.mongodb.com/manual/tutorial/query-documents/>

Understand the notion of MongoDB cursor and cursor timeout

Understand how query operators work

## PROGRAMMATICALLY

```
> db.collection.find({'name': 'gino'}) #MongoDB (see findOne)

cursor = db[collection].find({'name': 'gino'})
        #pymongo: cursor is iterable (find_one returns dict)

DBCursor cursor = coll.find(new BasicDBObject("name", "gino"));
        #Java: cursor.hasNext() and cursor.next()
```

## EXERCISE(S) 3.1

---

Define and execute the following queries

- Q1: Find videogames that can be played by 2 people
- Q2: Find title of boardgames that can be played by 2 people
- Q3: Find title and year of boardgames that can be played by 2 people sorted by year descending and title ascending
- Q4: Find titles videogames of category **Action**
- Q5: Find titles, year, and categories of videogames either of category **Action** or **Adventure**
- Q6: Find title and categories of videogames that are of type **Action AND Adventure**

## EXERCISE(S) 3.2

---

Define and execute the following queries

- Q7: Find title of boardgames published by Avalon Games and playable in no more than 120 min
- Q8: Find title of videogames in which the second category is Action
- Q9: Find title and votes of videogames that received at least one evaluation of 8 or more
- Q10: Find titles and platforms of videogames that are sold by Amazon at less than 40
- Q11: Find titles of videogames for which we have ratings
- Q12: Find titles of boardgames where minimum number of players is between 2 and 4
  - use `.explain("executionStats")` to examine results
  - try `db.boardgames.createIndex({min_players: 1})`
  - use `.explain("executionStats")` to examine results again

Download: [http://islab.di.unimi.it/~alfio/mdbexe/ginfo\\_games.tgz](http://islab.di.unimi.it/~alfio/mdbexe/ginfo_games.tgz)

```
mongorestore --db [your db name] [path/to/dump/folder/]
```

## AGGREGATION FRAMEWORK

---

Aggregations operations process data records and return computed results. Aggregation operations group values from multiple documents together, and can perform a variety of operations on the grouped data to return a single result.

## AGGREGATION PIPELINE

---

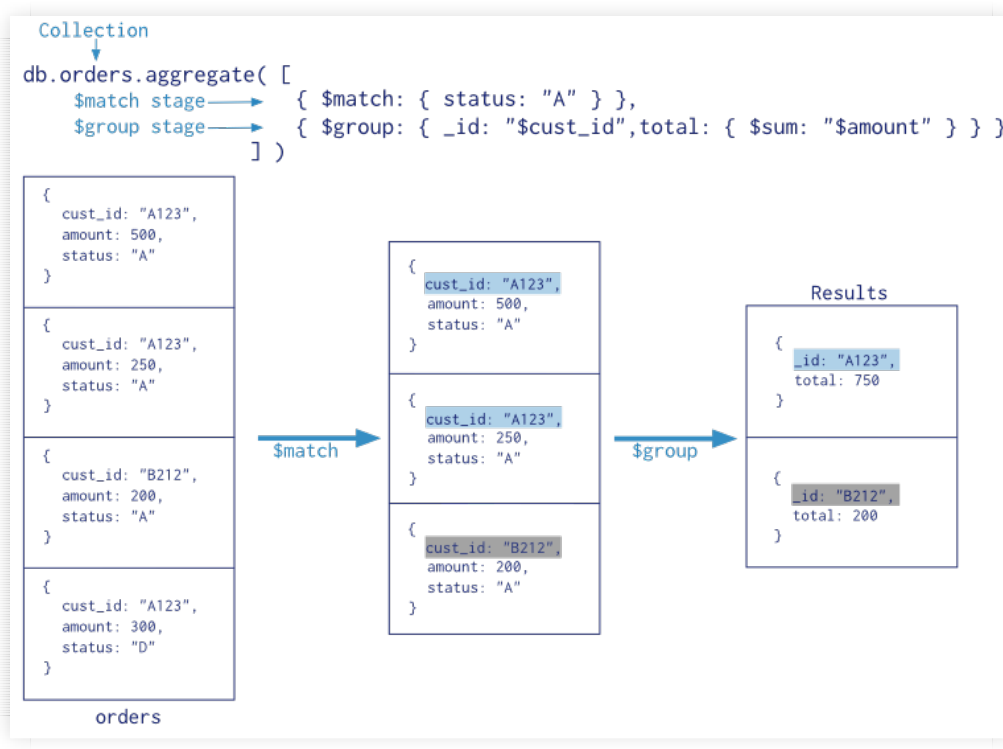
The aggregation pipeline is a **sequence of operations on data** where the only condition is that data returned from operation **i** are compatible with the input required from operation **i + 1**

Remember?

$$\pi_{A,B}(\sigma_{c=v}(R \bowtie_{x=y} S))$$

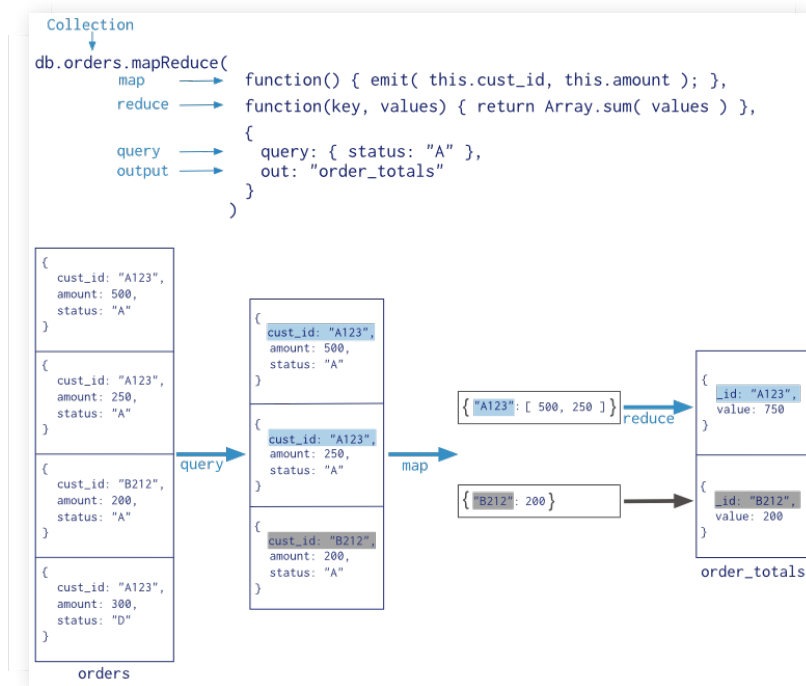
An aggregation pipeline has the following structure:

**db.collection.aggregate([O1, O2, ..., On])**



# ALTERNATIVE: MAPREDUCE

Map-reduce operations have two phases: a **map** stage that processes each document and **emits** one or more objects for each input document, and **reduce** phase that **combines** the output of the map operation



# \$GROUP

---

Groups input documents by a **specified identifier expression** and applies the **accumulator expression(s)**, if specified, to each group. Consumes all input documents and outputs one document per each distinct group. The output documents only contain the identifier field and, if specified, accumulated fields.

```
db.games.aggregate([
  {
    $group: {
      _id: "$type", //GROUP BY clause
      players: {
        "$avg": "$max_players" //ACCUMULATION (Aggregate operator)
      }
    }
  }
])
```

**Exercise:** find average number of minimum and maximum players per year



## SPECIAL CASE: ACCUMULATE ON ALL

---

```
db.games.aggregate([
  {
    $group: {
      _id: null, //GROUP BY clause
      players: {
        "$avg": "$max_players" //ACCUMULATION (Aggregate operator)
      }
      number: {"$sum": 1} // The way we implement COUNT
    }
  }
])
```

## MULTIPLE FIELD GROUPING

---

```
db.games.aggregate([
  {
    $group: {
      _id: {"year": "$year", "type": "$type"}, //GROUP BY clause
      players: {
        "$avg": "$max_players" //ACCUMULATION (Aggregate operator)
      }
    }
  }
])
```

# ACCUMULATORS

---

\$sum	Returns a sum of numerical values. Ignores non-numeric values.
\$avg	Returns an average of numerical values. Ignores non-numeric values.
\$first / \$last	Returns a value from the first/last document for each group. Order is only defined if the documents are in a defined order.
\$max / \$min	Returns the highest/lowest expression value for each group.
\$push / \$addToSet	Returns an array of expression values ( <b>unique for \$addToSet</b> ) for each group.
\$stdDevPop / \$stdDevSamp	Returns the population/sample standard deviation of the input values.

## \$MATCH

---

Filters the document stream to allow only matching documents to pass unmodified into the next pipeline stage. \$match uses standard MongoDB queries. For each input document, outputs either one document (a match) or zero documents (no match).

**Exercise:** calculate the number of **Action** games per year in the period 2010-2015

## \$PROJECT

---

Reshapes each document in the stream, such as by adding new fields or removing existing fields. For each input document, outputs one document.

```
db.games.aggregate([{$project: {publisher: "$publisher.name", year: 1}}])
```

## \$REDACT

---

Reshapes each document in the stream by restricting the content for each document based on information stored in the documents themselves. Incorporates the functionality of **\$project** and **\$match**. Can be used to implement field level redaction. For each input document, outputs either one or zero documents.

# \$UNWIND

---

Deconstructs an array field from the input documents to output a document for each element. Each output document replaces the array with an element value. For each input document, outputs  $n$  documents where  $n$  is the number of array elements and can be zero for an empty array.

**Exercise:** count games per category

## COMMENTED EXERCISE

---

Find average rating per month and category

```
u1 = {'$unwind': "$categories"}
u2 = {'$unwind': "$ratings"}
p = {'$project': {'month': {'$month': "$ratings.date"},
                  'category': "$categories", 'rate': "$ratings.vote"}}
g = {'$group': {'_id': {"month": "$month", "category": "$category"},
               "rate": {"$avg": "$rate"}}}
pipeline = [u1, u2, p, g]
db['games'].aggregate(pipeline)
```

## OTHER OPERATIONS

---

<code>\$sample</code> (3.2)	Randomly selects the specified number of documents from its input.
<code>\$lookup</code> (3.2)	Performs a left outer join to another collection in the same database to filter in documents from the “joined” collection for processing.
<code>\$out</code>	Writes the resulting documents of the aggregation pipeline to a collection. To use the <code>\$out</code> stage, it must be the last stage in the pipeline.

## EXERCISES

---

- A1: find years when more than 100 games have been published
- A2: find distribution of categories per year